

FACE RECOGNITION WITH DEEP LEARNING IN KERAS USING CNN

DESCRIPTION

Facial recognition is a biometric alternative that measures unique characteristics of a human face. Applications available today include flight check in, tagging friends and family members in photos, and “tailored” advertising. You are a computer vision engineer who needs to develop a face recognition programme with deep convolutional neural networks.

OBJECTIVE

Use a deep convolutional neural network to perform facial recognition using Keras.

DATASET DETAILS

ORL face database composed of 400 images of size 112 x 92. There are 40 people, 10 images per person. The images were taken at different times, lighting and facial expressions. The faces are in an upright position in frontal view, with a slight left-right rotation.

PRE-REQUISITES

Keras

Scikit Learn

STEPS TO BE FOLLOWED

1. Input the required libraries
2. Load the dataset after loading the dataset, you have to normalize every image.
3. Split the dataset
4. Transform the images to equal sizes to feed in CNN
5. Build a CNN model that has 3 main layers:
 - i. Convolutional Layer
 - ii. Pooling Layer
 - iii. Fully Connected Layer
6. Train the model
7. Plot the result
8. Iterate the model until the accuracy is above 90%

Step 1 : Input the required Libraries

```
In [28]: import keras
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
from tensorflow.keras.optimizers import Adam
from keras.callbacks import TensorBoard

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import accuracy_score
from keras.utils import np_utils
import itertools
```

Step 2 . Load the dataset after loading the dataset, you have to normalize every image.

```

In [29]: #Load dataset
data = np.load('ORL_faces.npz')

# Load the "Train Images"
x_train = data['trainX']

#normalize every image
x_train = np.array(x_train,dtype='float32')/255

x_test = data['testX']
x_test = np.array(x_test,dtype='float32')/255

# Load the Label of Images
y_train= data['trainY']
y_test= data['testY']

# show the train and test Data format
print('x_train : {}'.format(x_train[:]))
print('Y-train shape: {}'.format(y_train))
print('x_test shape: {}'.format(x_test.shape))

x_train : [[0.1882353  0.19215687 0.1764706  ... 0.18431373 0.18039216 0.180392
16]
[0.23529412 0.23529412 0.24313726 ... 0.1254902  0.13333334 0.13333334]
[0.15294118 0.17254902 0.20784314 ... 0.11372549 0.10196079 0.11372549]
...
[0.44705883 0.45882353 0.44705883 ... 0.38431373 0.3764706  0.38431373]
[0.4117647  0.4117647  0.41960785 ... 0.21176471 0.18431373 0.16078432]
[0.45490196 0.44705883 0.45882353 ... 0.37254903 0.39215687 0.39607844]]
Y-train shape: [ 0  0  0  0  0  0  0  0  0  0  0  0  0  1  1  1  1  1  1  1  1  1
1  1  1
 2  2  2  2  2  2  2  2  2  2  2  2  3  3  3  3  3  3  3  3  3  3
 4  4  4  4  4  4  4  4  4  4  4  4  5  5  5  5  5  5  5  5  5  5
 6  6  6  6  6  6  6  6  6  6  6  6  7  7  7  7  7  7  7  7  7  7
 8  8  8  8  8  8  8  8  8  8  8  8  9  9  9  9  9  9  9  9  9  9
10 10 10 10 10 10 10 10 10 10 10 10 11 11 11 11 11 11 11 11 11 11
12 12 12 12 12 12 12 12 12 12 12 12 13 13 13 13 13 13 13 13 13 13
14 14 14 14 14 14 14 14 14 14 14 14 15 15 15 15 15 15 15 15 15 15
16 16 16 16 16 16 16 16 16 16 16 16 17 17 17 17 17 17 17 17 17 17
18 18 18 18 18 18 18 18 18 18 18 18 19 19 19 19 19 19 19 19 19 19]
x_test shape: (160, 10304)

```

Step 3: Split the dataset

The dataset is splitted into two: Validation data and Train

Validation DataSet: this data set is used to minimize overfitting.If the accuracy over the training data set increases, but the accuracy over then validation data set stays the same or decreases, then you're overfitting your neural network and you should stop training.

Note: we usually use 30 percent of every dataset as the validation data but Here we only used 5 percent because the number of images in this dataset is very low.

```
In [30]: x_train, x_valid, y_train, y_valid= train_test_split(
        x_train, y_train, test_size=.05, random_state=1234,)
```

Step 4: Transform the images to equal sizes to feed in CNN

```
In [36]: im_rows=112
        im_cols=92
        batch_size=512
        im_shape=(im_rows, im_cols, 1)

        #change the size of images
        x_train = x_train.reshape(x_train.shape[0], *im_shape,)
        x_test = x_test.reshape(x_test.shape[0], *im_shape,)
        x_valid = x_valid.reshape(x_valid.shape[0], *im_shape,)

        print('x_train shape: {}'.format(y_train.shape[0]))
        print('x_test shape: {}'.format(y_test.shape))
```

```
x_train shape: 228
x_test shape: (160,)
```

Step 5: Build CNN model: CNN have 3 main layer:

- 1-Convolutional layer
- 2- pooling layer
- 3- fully connected layer

we could build a new architecture of CNN by changing the number and position of layers.

In [37]: *#filters= the depth of output image or kernels*

```
cnn_model= Sequential([
    Conv2D(filters=36, kernel_size=7, activation='relu', input_shape= im_shape),
    MaxPooling2D(pool_size=2),
    Conv2D(filters=54, kernel_size=5, activation='relu', input_shape= im_shape),
    MaxPooling2D(pool_size=2),
    Flatten(),
    Dense(2024, activation='relu'),
    Dropout(0.5),
    Dense(1024, activation='relu'),
    Dropout(0.5),
    Dense(512, activation='relu'),
    Dropout(0.5),
    #20 is the number of outputs
    Dense(20, activation='softmax')
])

cnn_model.compile(
    loss='sparse_categorical_crossentropy',#'categorical_crossentropy',
    optimizer=Adam(learning_rate=0.0001),
    metrics=['accuracy']
)
```

```
In [38]: #Display the model summary
cnn_model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_6 (Conv2D)	(None, 106, 86, 36)	1800
max_pooling2d_6 (MaxPooling2)	(None, 53, 43, 36)	0
conv2d_7 (Conv2D)	(None, 49, 39, 54)	48654
max_pooling2d_7 (MaxPooling2)	(None, 24, 19, 54)	0
flatten_3 (Flatten)	(None, 24624)	0
dense_12 (Dense)	(None, 2024)	49841000
dropout_9 (Dropout)	(None, 2024)	0
dense_13 (Dense)	(None, 1024)	2073600
dropout_10 (Dropout)	(None, 1024)	0
dense_14 (Dense)	(None, 512)	524800
dropout_11 (Dropout)	(None, 512)	0
dense_15 (Dense)	(None, 20)	10260
=====		
Total params: 52,500,114		
Trainable params: 52,500,114		
Non-trainable params: 0		

Step 6: Train the model

```
In [39]: history=cnn_model.fit(np.array(x_train), np.array(y_train), batch_size=512,epochs=250,
validation_data=(np.array(x_valid),np.array(y_valid)))
```

```
Epoch 1/250
1/1 - 28s - loss: 3.0042 - accuracy: 0.0482 - val_loss: 2.9866 - val_accuracy: 0.0000e+00
Epoch 2/250
1/1 - 1s - loss: 3.0220 - accuracy: 0.0482 - val_loss: 3.0040 - val_accuracy: 0.0000e+00
Epoch 3/250
1/1 - 1s - loss: 3.0049 - accuracy: 0.0263 - val_loss: 3.0075 - val_accuracy: 0.0000e+00
Epoch 4/250
1/1 - 1s - loss: 3.0315 - accuracy: 0.0570 - val_loss: 3.0117 - val_accuracy: 0.0000e+00
Epoch 5/250
1/1 - 1s - loss: 2.9906 - accuracy: 0.0614 - val_loss: 3.0124 - val_accuracy: 0.0000e+00
Epoch 6/250
1/1 - 1s - loss: 2.9994 - accuracy: 0.0351 - val_loss: 3.0121 - val_accuracy: 0.0000e+00
Epoch 7/250
1/1 - 1s - loss: 2.9971 - accuracy: 0.0500 - val_loss: 3.0111 - val_accuracy: 0.0000e+00
```

```
In [40]: scor = cnn_model.evaluate( np.array(x_test), np.array(y_test), verbose=0)



print('test los {:.4f}'.format(scor[0]))
print('test acc {:.4f}'.format(scor[1]))
```

```
test los 0.3114
test acc 0.9500
```

Step 7: Plot the Results

```
In [44]: # list all data in history
print(history.history.keys())
# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```



Step 8: Iterate the model until the accuracy is above 90%

```
In [54]: predicted =np.array( cnn_model.predict(x_test))
          #print(predicted)
          #print(y_test)

          predict_x=cnn_model.predict(x_test)
          ynew=np.argmax(predict_x,axis=1)

          #ynew = cnn_model.predict_classes(x_test)

          Acc=accuracy_score(y_test, ynew)
          print("accuracy : ",Acc*100,"%")
```

accuracy : 95.0 %

```
In [55]: #Confusion Matrix
```

```
In [57]: cnf_matrix=confusion_matrix(np.array(y_test), ynew)

y_test1 = np_utils.to_categorical(y_test, 20)

def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        #print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    #print(cm)
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.show()
```


```
In [58]: #for classes 0-9
print('Confusion matrix, without normalization')
print(cnf_matrix)

plt.figure()
plot_confusion_matrix(cnf_matrix[1:10,1:10], classes=[0,1,2,3,4,5,6,7,8,9],
                      title='Confusion matrix, without normalization')
```

Confusion matrix, without normalization

```
[[8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 0 2 0 0]
 [0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 2 0 6 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0]
 [2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 8 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 8 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 2 0 0 0 0 0 0 6 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 8 0]]
```


Confusion matrix, without normalization



```
In [59]: #for classes 10-19
plt.figure()
plot_confusion_matrix(cnf_matrix[11:20,11:20], classes=[10,11,12,13,14,15,16,17,18,19],
                      title='Confusion matrix, without normalization')

print("Confusion matrix:\n%s" % confusion_matrix(np.array(y_test), ynew))
print(classification_report(np.array(y_test), ynew))
```

Confusion matrix, without normalization



Confusion matrix:

```
[8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 6 0 0 0 0 0 0 0 0 0 0 0 0 2 0 0]
[0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 2 0 6 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 8 0 0 0 0 0 0 0]
[2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 6 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 8 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 8 0 0]
[0 0 0 0 0 0 0 0 0 0 2 0 0 0 0 0 0 0 6 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 8]]
```

precision recall f1-score support

	0	0.80	1.00	0.89	8
	1	1.00	1.00	1.00	8
	2	1.00	1.00	1.00	8
	3	1.00	1.00	1.00	8
	4	1.00	0.75	0.86	8
	5	1.00	1.00	1.00	8
	6	1.00	1.00	1.00	8
	7	0.80	1.00	0.89	8
	8	1.00	1.00	1.00	8
	9	1.00	0.75	0.86	8
	10	0.80	1.00	0.89	8
	11	1.00	1.00	1.00	8
	12	1.00	1.00	1.00	8
	13	1.00	1.00	1.00	8
	14	1.00	1.00	1.00	8
	15	1.00	0.75	0.86	8
	16	1.00	1.00	1.00	8
	17	0.80	1.00	0.89	8
	18	1.00	0.75	0.86	8
	19	1.00	1.00	1.00	8
	accuracy			0.95	160
	macro avg	0.96	0.95	0.95	160
	weighted avg	0.96	0.95	0.95	160



In []:

